

NLTK WordNet Exploration

Author: Andrew Sen

Date: 9/25/2022

WordNet is a hierarchical organization of nouns, verbs, adjectives, and adverbs. For each word, WordNet includes information about definitions of the word, synonym sets called synsets, usage examples, and relations to other words.

Synsets for Nouns

I will use the noun "game" and explore its associated synsets with WordNet. First, I will output all synsets for "game."

```
In [ ]: from nltk.corpus import wordnet as wn

wn.synsets('game')
```

```
Out[ ]: [Synset('game.n.01'),
         Synset('game.n.02'),
         Synset('game.n.03'),
         Synset('game.n.04'),
         Synset('game.n.05'),
         Synset('game.n.06'),
         Synset('game.n.07'),
         Synset('plot.n.01'),
         Synset('game.n.09'),
         Synset('game.n.10'),
         Synset('game.n.11'),
         Synset('bet_on.v.01'),
         Synset('crippled.s.01'),
         Synset('game.s.02')]
```

Using the first synset in the list, I will output its definition, usage examples, and lemmas.

```
In [ ]: game = wn.synset('game.n.01')
# extracting definitions
print(f'Definition:\n{game.definition()}\n')
# extracting usage examples
print(f'Usage Examples:\n{game.examples()}\n')
# extracting lemmas
print(f'Lemmas:\n{game.lemmas()}\n')
```

Definition:

a contest with rules to determine a winner

Usage Examples:

['you need four people to play this game']

Lemmas:

[Lemma('game.n.01.game')]

We can also use WordNet to traverse up a word's hierarchy of hypernyms.

```
In [ ]: # extracting hypernyms of 'game'
hyp = game.hypernyms()[0]
while hyp:
    print(hyp)
    if hyp.hypernyms():
        hyp = hyp.hypernyms()[0]
    else:
        break
```

Synset('activity.n.01')

Synset('act.n.02')

Synset('event.n.01')

Synset('psychological_feature.n.01')

Synset('abstraction.n.06')

Synset('entity.n.01')

As we can see with this example, WordNet has a topmost hypernym for nouns in the form of the 'entity' synset. Under this system, all other nouns are a hyponym of 'entity.' As we will see, this is unlike how verbs are organized in WordNet.

Synsets for Verbs

As we did for the noun 'game,' we will do the same to explore the synsets of the verb 'play.'

```
In [ ]: wn.synsets('play')
```

```

Out[ ]: [Synset('play.n.01'),
        Synset('play.n.02'),
        Synset('play.n.03'),
        Synset('maneuver.n.03'),
        Synset('play.n.05'),
        Synset('play.n.06'),
        Synset('bid.n.02'),
        Synset('play.n.08'),
        Synset('playing_period.n.01'),
        Synset('free_rein.n.01'),
        Synset('shimmer.n.01'),
        Synset('fun.n.02'),
        Synset('looseness.n.05'),
        Synset('play.n.14'),
        Synset('turn.n.03'),
        Synset('gambling.n.01'),
        Synset('play.n.17'),
        Synset('play.v.01'),
        Synset('play.v.02'),
        Synset('play.v.03'),
        Synset('act.v.03'),
        Synset('play.v.05'),
        Synset('play.v.06'),
        Synset('play.v.07'),
        Synset('act.v.05'),
        Synset('play.v.09'),
        Synset('play.v.10'),
        Synset('play.v.11'),
        Synset('play.v.12'),
        Synset('play.v.13'),
        Synset('play.v.14'),
        Synset('play.v.15'),
        Synset('play.v.16'),
        Synset('play.v.17'),
        Synset('play.v.18'),
        Synset('toy.v.02'),
        Synset('play.v.20'),
        Synset('dally.v.04'),
        Synset('play.v.22'),
        Synset('dally.v.01'),
        Synset('play.v.24'),
        Synset('act.v.10'),
        Synset('play.v.26'),
        Synset('bring.v.03'),
        Synset('play.v.28'),
        Synset('play.v.29'),
        Synset('bet.v.02'),
        Synset('play.v.31'),
        Synset('play.v.32'),
        Synset('play.v.33'),
        Synset('meet.v.10'),
        Synset('play.v.35')]

```

I will select the first verb synset in the list and extract its definition, usage examples, and lemmas.

```
In [ ]: play = wn.synset('play.v.01')
# extracting definitions
print(f'Definition:\n{play.definition()}\n')
# extracting usage examples
print(f'Usage Examples:\n{play.examples()}\n')
# extracting Lemmas
print(f'Lemmas:\n{play.lemmas()}\n')
```

Definition:
participate in games or sport

Usage Examples:
['We played hockey all afternoon', 'play cards', 'Pele played for the Brazilian teams in many important matches']

Lemmas:
[Lemma('play.v.01.play')]

We can also traverse the heirarchy of hypernyms for a given verb.

```
In [ ]: # extracting hypernyms of 'play'
hyp = play.hypernyms()[0]
while hyp:
    print(hyp)
    if hyp.hypernyms():
        hyp2 = hyp.hypernyms()[0]
    else:
        break
```

Synset('compete.v.01')

Here we see that the topmost hypernym for 'play' is 'compete.' This cannot be a general hypernym for all verbs. This shows that, unlike with nouns, WordNet does not categorize all verbs as being hyponyms to some universal umbrella verb.

The `morphy()` function returns the base form of a word. We can use it to confirm that certain words are just different forms of the word 'play.'

```
In [ ]: forms = ['played', 'playing', 'plays', 'player']
print("word: base\n-----")
for f in forms:
    print(f'{f}: {wn.morphy(f)}')
```

word: base

played: play
playing: playing
plays: play
player: player

Similarity and Word Sense Disambiguation

WordNet includes the Wu-Palmer algorithm for determining the similarity between two words. NLTK also has an implementation of the Lesk algorithm for determining which definition of a word is being used in a given sentence. To test both, I will use two different forms of the word 'punch.'

First, I will output the definitions of each synset of 'punch' to pick the right synsets to work with.

```
In [ ]: for ss in wn.synsets('punch'):
        print(ss.name() + ': ' + ss.definition())

# Lesk

punch.n.01: (boxing) a blow with the fist
punch.n.02: an iced mixed drink usually containing alcohol and prepared for multiple servings; normally served in a punch bowl
punch.n.03: a tool for making holes or indentations
punch.v.01: deliver a quick blow to
punch.v.02: drive forcibly as if by a punch
punch.v.03: make a hole into or between, as for ease of separation
```

I will use `punch.n.02` and `punch.v.01` for this example. Now let's use Wu-Palmer to determine their similarity.

```
In [ ]: punch_hit = wn.synset('punch.v.01')
        punch_drink = wn.synset('punch.n.02')

        wn.wup_similarity(punch_hit, punch_drink)
```

```
Out[ ]: 0.13333333333333333
```

We see that Wu-Palmer gave the two senses of 'punch' a low similarity score, which is to be expected.

Now let's use Lesk to disambiguate the use of 'punch' in a sentence.

```
In [ ]: from nltk.wsd import lesk

sentence = ['I', 'walked', 'to', 'the', 'table',
            'and', 'grabbed', 'some', 'punch']
print(lesk(sentence, 'punch'))

sentence2 = ['I', 'wanted', 'to', 'punch', 'him']
print(lesk(sentence2, 'punch'))
```

```
Synset('punch.n.02')
Synset('punch.v.02')
```

Here we see that Lesk was able to correctly disambiguate the first sentence. The disambiguation suggested for the second sentence was not what I expected, but the definitions of `punch.v.01` and `punch.v.02` are similar nonetheless.

SentiWordNet

SentiWordNet is a tool for programmatically determining the sentiment of a piece of text. Given some text, it will assign scores in positivity, negativity, and objectivity.

For this example, I will use the word 'attack.'

```
In [ ]: from nltk.corpus import sentiwordnet as swn

for ss in swn.senti_synsets('attack'):
    print(ss)
```

```
<attack.n.01: PosScore=0.0 NegScore=0.0>
<attack.n.02: PosScore=0.0 NegScore=0.0>
<fire.n.09: PosScore=0.125 NegScore=0.5>
<approach.n.01: PosScore=0.0 NegScore=0.0>
<attack.n.05: PosScore=0.0 NegScore=0.0>
<attack.n.06: PosScore=0.0 NegScore=0.0>
<attack.n.07: PosScore=0.0 NegScore=0.25>
<attack.n.08: PosScore=0.0 NegScore=0.125>
<attack.n.09: PosScore=0.25 NegScore=0.125>
<attack.v.01: PosScore=0.0 NegScore=0.0>
<attack.v.02: PosScore=0.0 NegScore=0.0>
<attack.v.03: PosScore=0.0 NegScore=0.5>
<assail.v.01: PosScore=0.0 NegScore=0.375>
<attack.v.05: PosScore=0.0 NegScore=0.0>
<attack.v.06: PosScore=0.0 NegScore=0.0>
```

Now I will make up a sentence and find the polarity of each word in the sentence.

```
In [ ]: sentence = 'I really hate drinking iced coffee'
tokens = sentence.split()

for token in tokens:
    ss = lesk(tokens, token) # use Lesk to get best synset
    print(f'{token}: {swn.senti_synset(ss.name())}')
```

```
I: <one.s.01: PosScore=0.0 NegScore=0.25>
really: <very.r.01: PosScore=0.25 NegScore=0.25>
hate: <hate.v.01: PosScore=0.0 NegScore=0.75>
drinking: <drink.n.02: PosScore=0.0 NegScore=0.0>
iced: <ice.v.03: PosScore=0.0 NegScore=0.0>
coffee: <coffee_bean.n.01: PosScore=0.0 NegScore=0.0>
```

We see that the word 'I' interestingly has a slightly negative polarity. The word 'really' has equal positive and negative polarities, which makes sense because it can precede either a positive or negative word. 'Hate' is very negative as to be expected, and the remaining words have no polarity.

In a real NLP application, it would be useful to have these sentiment scores because it gives extra information about the meaning of a text. Sentiment information could be used as extra factors in a given language model.

Collocations

A collocation is a group of words that, when put together, refer to a particular thing or action, and the same effect is not achieved if one of the words is replaced with a synonym. An example would be the term 'fast food.' 'Fast food' means something very particular, and saying 'quick food' either sounds wrong or is referring to a different concept altogether.

We will list the collocations found in one of NLTK's built-in texts.

```
In [ ]: from nltk.book import text4
        text4.collocations()
```

```
United States; fellow citizens; years ago; four years; Federal
Government; General Government; American people; Vice President; God
bless; Chief Justice; one another; fellow Americans; Old World;
Almighty God; Fellow citizens; Chief Magistrate; every citizen; Indian
tribes; public debt; foreign nations
```

I will now calculate the mutual information of the collocation 'Federal Government.'

```
In [ ]: import math

text = ' '.join(text4.tokens)
vocab = len(set(text))
fg = text.count('Federal Government') / vocab
print("p(Federal Government) = ", fg)
f = text.count('Federal') / vocab
print("p(Federal) = ", f)
g = text.count('Government') / vocab
print('p(Government) = ', g)
pmi = math.log2(fg / (f * g))
print('pmi = ', pmi)
```

```
p(Federal Government) = 0.38095238095238093
p(Federal) = 0.7738095238095238
p(Government) = 4.023809523809524
pmi = -3.0309298265318785
```

A negative pmi indicates that 'Federal Government' is not likely to be a collocation in this text. Since NLTK did consider it a collocation, we can assume that NLTK uses some other means of determining whether or not a phrase is a collocation.

